

# Mark – Your Personal AI Investment Assistant

Ciro Francesco Amabile, Vincenzo Carbone, Gerardo D'Arco

# Introduction

Our goal is to create a virtual financial assistant capable of answering questions about publicly traded companies.

To ensure that the data is always up to date, we plan to use a **VPS (Virtual Private Server)**, where the pipeline will run periodically using crontab.

Mark is powered by a **RAG (Retrieval-Augmented Generation)** pipeline that combines:

- **Retrieval:** the system searches the financial data for the most relevant content related to the user's question
- **Augmented Generation:** the retrieved content is used as context to generate an accurate answer with an advanced language model (GPT)

In summary, Mark combines:

- Collected and cleaned financial data
- NLP techniques for semantic understanding
- GPT models for clear and professional answers

# Project Structure

```
Directory:
mark-ai/
├── README.md
├── requirements.txt
├── .env
├── /data/
│   ├── /raw/fundamentals/
│   ├── /raw/market_data/
│   ├── /pdf_raw/
│   ├── /text_clean/
│   ├── /chunks/
│   └── /index/
├── /scripts/
│   ├── download_fundamentals.py
│   ├── download_market_data.py
│   ├── download_pdfs.py
│   ├── extract_text.py
│   ├── clean_text.py
│   ├── chunk_financials.py
│   ├── embed_and_index.py
│   ├── retriever.py
│   ├── llm_wrapper.py
│   └── ask_mark.py
└── /web/
    └── mark_streamlit_app.py
```

# Project Structure - In Depth

## Main directory: mark-ai/

- data/

Contains all raw and processed data. Subdivided into:

- raw/fundamentals, raw/market\_data: data downloaded via API
- pdf\_raw: company PDFs obtained via web scraping
- text\_clean: cleaned and formatted text extracted from PDFs
- chunks: information blocks ready for embedding
- index: FAISS vector database and mapping

- scripts/

Contains all Python scripts for the pipeline. Each file performs a specific function: scraping, cleaning, embedding, retrieval, and response generation.

- web/

Future upgrade for a web-based user interface.

- README.md, requirements.txt, .env

Support files: project documentation, required libraries, and environment variables (e.g., API keys).

# Data: Collection and Cleaning

## 1. Automated data collection:

- We use yfinance to obtain:
  - **Fundamental data:** revenue, net income, P/E ratio, operating margin, ROE, etc.
  - **Market data:** current price, historical prices, beta, trading volumes
- Outputs are saved as .json or .csv files in the data/raw/ directory
- In parallel, we can download **PDFs** (quarterly reports, ESG documents, press releases) via web scraping from official websites (Investor Relations, SEC, CONSOB)

## 2. Cleaning and pre-processing:

- Conversion of numbers and percentages into standardized formats
- Removal of noise from text (page numbers, repeated headers, unwanted symbols)
- Normalization of content to ensure consistency across different sources
- Production of "clean" text ready for chunk segmentation

# Scripts - Detailed Overview

## 1. Data Collection

- `download_fundamentals.py`: downloads key variables (revenue, income, P/E, etc.)
- `download_market_data.py`: retrieves stock prices, trading volumes, beta, etc

## 2. PDFs and Text

- `download_pdfs.py`: scrapes PDFs from official company websites
- `extract_text.py`: extracts raw text from the downloaded PDFs
- `clean_text.py`: cleans and normalizes the extracted content

## 3. Chunk Preparation

- `chunk_financials.py`: segments the text into semantic blocks ("chunks") with metadata

## 4. Embedding and Indexing

- `embed_and_index.py`: generates numerical embeddings and creates a FAISS index

## 5. User Interaction

- `retriever.py`: retrieves the most relevant chunks for the given question
- `llm_wrapper.py`: builds the prompt and calls the ChatGPT API
- `ask_mark.py`: CLI script that allows direct interaction with Mark

# Embedding and FAISS

## 1. Semantic Embedding

- Each **text chunk** (300–500 words) is transformed into a high-dimensional **numerical vector** that captures its meaning.
- We use OpenAI's text-embedding-ada-002 model, one of the most efficient and semantically accurate models available.
- Result: each information block becomes a set of numbers that reflects its semantic essence.

## 2. FAISS (Facebook AI Similarity Search)

- The vectors are stored and indexed in a **vector store** using the FAISS library, designed for fast similarity search across millions of vectors.
- When a user submits a question, it is also embedded into a vector.
- The system compares this vector with all chunk vectors and retrieves the most semantically similar ones.

# Final expected Output

## **Example user question:**

"What is Apple's P/E ratio and how does it compare to Microsoft's?"

### **1. Retrieval of relevant chunks**

- The system converts the question into a vector embedding.
- Using FAISS, it retrieves the most semantically similar chunks (e.g., data on Apple and Microsoft).

### **2. Prompt construction**

- The selected chunks are formatted as "context" for GPT.
- A structured prompt is built, including both the relevant data and the user's question.

### **3. Answer generation**

- The GPT-3.5/4 model analyzes the information and generates a clear, professional answer.
- Output: a comparison of the P/E ratios, with possible comments on high/low values and market implications.



# Scalability: Future Extensions

The project is designed to be easily scalable. Potential extensions include:

## 1. Integration with new data sources

- Expanding data sources via **APIs** (e.g., Bloomberg)
- Automated **web scraping** to collect data from financial portals, Investor Relations pages, and open-access databases

## 2. Market sentiment analysis

- Scraping of news and articles to perform **sentiment analysis**, useful for anticipating market reactions to events or announcements

## 3. Regulatory analysis and legal transparency

- Extraction of content from public institutional sources to generate responses that also reflect the **rights and obligations** of a fully informed investor

## 4. New user interfaces

- Integration with a **Telegram bot** for quick queries
- Development of a **web dashboard** using Streamlit or FastAPI

# Thank you for your attention!

Project developed by:

- **Ciro Francesco Amabile**
- **Vincenzo Carbone**
- **Gerardo D'Arco**

*“Investing is no longer just about numbers — it's about automated intelligence.”*